

What should I notice? Evaluating memorability of events for abductive inference.

*† Étienne Houzé, * Jean-Louis Dessalles, * Ada Diaconescu, † David Menga

* Télécom Paris, IP Paris, *Palaiseau*, France

email: {first}.{second}@telecom-paris.fr

† EDF R&D, *Palaiseau*, France

email: {first}.{second}@edf.fr

Abstract—When confronted to an unprecedented situation, humans typically show good performance in quickly identifying noticeable past events and proposing them as possible causal hypotheses. This kind of abductive inference is widely overlooked in modern AI approaches which rely on massive datasets to learn associative patterns. Our proposal is to formalize and compute a “memorability” score over a memory of various recorded events from a cyber-physical system. This score can later be used either to select only relevant information to be remembered, or to propose causal hypotheses in unusual situations, on demand. As such, the approach aims at being complementary to more traditional learning-focused techniques. We provide theoretical ground for our approach by using and extending existing results and ideas from Algorithmic Information Theory and provide an implementation example, showing practical results in a smart-home scenario.

Index Terms—Complexity, Algorithmic Information Theory, Simplicity, Abduction, Surprise

I. INTRODUCTION

As a user has just turned on the TV in her all-equipped living-room, the lights dim and the window blinds lower. Intrigued by this behavior, she quickly infers that both light dimming and the blinds closing occurred as a consequence of the TV being turned on. How did she come to this conclusion? By performing *abductive inference* [1], which is a key element of humans’ ability to understand the world: from the observed consequences, infer the possible causes.

In this example, there are mainly three possibilities to come to the conclusion. If the user knows how the smart living-room system works, if she knows the set of applied rules or parameters, she has access to a causal knowledge she can use for the abduction. Otherwise, if she has no knowledge about the system but the behavior occurs frequently, examining past correlations will reveal that the television frequently causes the blinds to close and the lights to dim. Else, if this is the first instance of the TV being turned on in this living-room, there is no previous occurrence upon which the user may rely to draw conclusions. However, she remains able to infer the TV as a possible cause for the observed reaction: by noticing the TV as a memorable recent event (since it is its first occurrence), she proposes it as a possible cause for the questioned phenomenon. This example shows how humans are able to use distinct methods to perform abductive tasks and infer new knowledge. While abduction of

the first two kinds can be automated using knowledge bases and statistical approaches, this paper proposes an approach to evaluate the memorability of events, allowing machines to perform “memorability-based” abduction.

The difficulty of implementing this kind of abduction for machines comes from several factors. First, events can be of different nature, and not directly comparable. In fact, even for specific systems such as smart homes, events range from device removal to a presence detection or an unusually high temperature. How can one then assess which ones are more memorable? Furthermore, even for comparable events, different characteristics can be put forward to argue for the most memorable: is a record-high temperature 47 days ago more memorable than the small deviation recorded just 3 minutes ago? To this day, no current system proposes to use all this information from various event types from different devices to compute a unified metrics of “memorability”.

To tackle this challenge, we rely on the following observation: while all events, regardless of their characteristics or nature, can be uniquely described using a combination of qualifiers, the most memorable ones are likely to require less words to be described. Think, for instance, of “last year’s hottest day” and “the 182nd day from 7 years ago”: the former seems simpler to describe. To quantify the notion, we can evaluate the complexity of each description, taking into account both the complexity of concept words (a date of occurrence, a temperature ranking), and the arguments (the hottest, the 182nd, 7). The resulting value defines the *description complexity* of the events. As our intuition is that memorable events require simpler and less numerous qualifiers to be unambiguously described than other unremarkable ones, they will stand out as simpler in terms of description complexity.

For machines to implement and compute description complexity, we need a formal canvas and computation methods which coincide with human intuition. To do so, we rely on Algorithmic Information Theory, which also appears to be consistent with the human perception of complexity [2]–[4]. Our method is as follows: we consider events as being unordered elements from a base *memory*. To reproduce the language features applicable to events, we use *predicates*, i.e. functions assigning a boolean value to an event and eventual arguments. For instance, the predicate `date(·, 1_year)` assigns

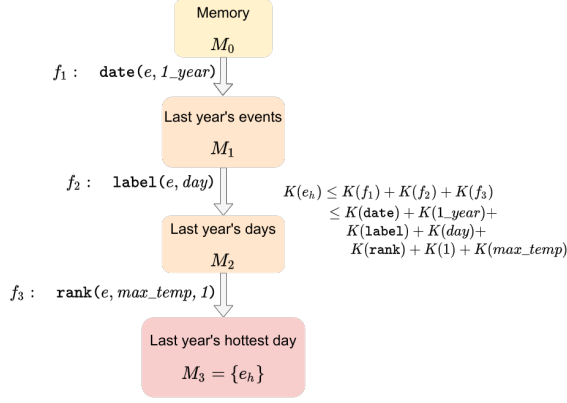


Fig. 1. Retrieving an event through successive predicative filters. From the base memory (yellow), successive filters select events satisfying the associated predicate (grey arrows). For example, filter f_1 selects events from last year, i.e. which satisfy the predicate $\text{date}(\text{event}, 1_year)$. In this case, successively applying filters f_1 , f_2 and f_3 yields a unique event e_h , last year's hottest day. The complexity of this event can then be upper bounded by the complexity of the three filters, since they give a unambiguous way to describe the event in the memory.

true to events that occurred last year. Selecting all events, from the memory, that satisfy a given predicate corresponds to a *filter* operation, and yields another memory, a subset of the former. This filtering operation can then be repeated, each iteration, selecting fewer events, until a singleton memory is reached, meaning that the sequence of predicates can *retrieve* unambiguously the unique remaining event. The description complexity of this event can thus be upper-bounded by the number of bits required to describe the filters used in the retrieval process. Figure 1 illustrates this evaluation for the event e_h : “last year’s hottest day”.

The rest of this paper is organized as follows: we first briefly introduce some notions of Algorithmic Information Theory and its applications in section II, then we explicit our methods to compute the memorability of past events in cyber-physical systems III. We illustrate this approach by providing an implementation relying on a few predicates on events from a smart home simulation (section IV). In section V we briefly review other related works using complexity theory or trying to explain smart homes and how our work can be linked with them before exploring potential extensions of our work in section VI

II. THEORETICAL BACKGROUND

Kolmogorov complexity formally quantifies the amount of information required for the computation of a finite binary string¹ (or any object represented by a finite binary string) [2], [5]. For such a binary string s , its complexity $K(s)$, is

¹While the definition still holds for some infinite binary strings (think of the representation of the decimals of π), we restrict ourselves to finite strings in this paper.

the bit length of the shortest program p which, if given as input to a universal Turing Machine U , outputs s .

$$K_U(s) = \min_p \{l(p) | U(p) = s\} \quad (1)$$

The first notable property of this definition is its universality: while the choice of the Turing machine U used for the computations appears in the definition of eq. 1, all results stand, up to an additional constant. Think, for instance, how any Turing-complete programming language can be turned into any other language, using a compiler program. In fact, since any Turing machine U' can be emulated by U from a finite program p_U , we have the following inequality:

$$K_U(s) \leq l(p_U) + K_{U'}(s) \quad (2)$$

From this first result, we can then define a universal complexity $K(s)$, which no longer depends on the choice of a Turing machine U , such that, for any machine U ,

$$\forall U \in \text{TM}, \forall s, |K(s) - K_U(s)| \leq C_U \quad (3)$$

where the additional constant C_U does not depend on the object s .

It is important to note that the notion of Kolmogorov complexity is not related to computation complexity, as there is no requirement on the execution time of the programs, only their length in bits matters for the computation of complexity. In fact, it can be shown that Kolmogorov complexity is not computable [2]. The sketch of the proof is the following: since the definition of complexity from Equation. 1 requires to find the shortest of all programs on a given Turing machine outputting s , it requires to run all programs up to a given length and compare their result to s . However, being able to do so means that we would be able to tell whether these programs terminate. Since the termination of programs is famously undecidable, by extension, the computation of Kolmogorov complexity is impossible. One can however easily approximate it with upper bounds, by exhibiting a program outputting s .

Interestingly, Kolmogorov’s definition of complexity matches the intuitive notion and perception of complexity from a human standpoint. For instance, the complexity of short binary strings evaluated in [4] shows similar results to human perception of complex strings and patterns. More recently, [6] used Kolmogorov complexity to solve analogies and showed results close to human expectations.

This bridge between Algorithmic Information Theory and human perception of complexity is extended to define the notion of simplicity and unexpectedness, which are considered of uttermost importance in cognitive science [7]. [3] proposes a formal definition of the unexpectedness $U(e)$ of an event, as the difference between an a-priori expected causal complexity $C_w(e)$ and the actual observed complexity $C(e)$.

$$U(e) = C_w(e) - C(e) \quad (4)$$

This result comes from the understanding that, while Kolmogorov complexity is computed with Turing machine, it can be used as an elegant proxy for modeling information

processing in the human brain, and thus helps design a notion of simplicity or complexity of events.

The definition of Equation 4 allows to model phenomenons such as coincidences: imagine that you happen to run into someone in a park. If this person has no particular link to you, the event will be quite trivial: the complexity of describing this person will be equivalent to distinguishing it from the global population, which will be roughly equivalent as the complexity of describing that this person happens to be in the same park as you. On the other hand, if you run into your best friend in a park, as the complexity of describing your best friend is significantly lower, the description complexity $C(e)$ drops while the causal complexity remains $C_W(e)$ unchanged. Therefore the event becomes unexpected.

Using these insights from AIT, we define the memorability $M(e)$ of an event as the absolute difference between the description complexity $K_d(e)$ of an event and its expected description complexity $K_{exp}(e)$:

$$M(e) = |K_{exp}(e) - K_d(e)| \quad (5)$$

Contrary to the definition of unexpectedness from Equation 4, we use an absolute value: this is done to acknowledge events more complex than expected as memorable². In the next section, we formally define the description complexity K_d and the expected complexity K_{exp} of events and detail how we can compute approximations of these quantities.

III. COMPUTING THE MEMORABILITY OF EVENTS

A. Retrieving an event

Computing the memorability of events begins with the introduction of formal definitions of events, how they are stored, how we can describe them and how these descriptions can be used to retrieve them.

We define *events* as data points augmented with a *label* indicating their nature (temperature event, failure event, addition/removal of a device) and a timestamp of occurrence. Formally:

$$e = (l, t, \mathcal{D}) \quad (6)$$

where l is the label, t the timestamp and \mathcal{D} a multi-dimensional data point representing the various characteristics of e : its duration, the maximum temperature reached, the sensor name, its position, etc. Labels can also be considered as classes of events, of which each event is a particular instance.

To model how humans are able to describe events by using qualifiers, we use *predicates*: function operating on events, eventually taking additional parameters, returning a boolean value: $\pi(e, a_1, a_2, \dots, a_n) \mapsto \{0, 1\}$ is a predicate of arity n operating on event e . In the rest of this paper, we will prefer the equivalent notation $\pi(e, k) \mapsto \{0, 1\}$, where k is a binary string encoding the sequence of arguments a_1, \dots, a_n . Using

this notation, the predicate π becomes a boolean function operating on $\mathbf{E} \times \{0, 1\}^*$:

$$\pi : \begin{cases} \mathcal{M} \times \{0, 1\}^* & \mapsto \{0, 1\} \\ (e, k) & \mapsto \pi_k(e) \end{cases} \quad (7)$$

An example of predicate is to take $\pi = \text{year}$ and k a string encoding the number 1, thus constructing the predicate $\text{year}(e, 1)$, which tells whether the event e occurred 1 years ago.

As events occur, they are stored in a *memory* M_0 . As they are not directly comparable, the memory M_0 can be considered as having the structure of an unordered set. We denote by \mathcal{M} the subsets of M_0 . By extension, elements of \mathcal{M} , i.e. subsets of M_0 , are also called *memories*.

By applying a given predicate π_k to all events contained in a memory $M \subseteq M_0$, and selecting only events satisfying π_k yields another memory $M_1 \subseteq M \subseteq M_0$. We call this operation a *filter*:

$$f_{\pi, k} : \begin{cases} \mathcal{M} & \mapsto \mathcal{M} \\ M & \mapsto \{e \in M | \pi_k(e)\} \end{cases} \quad (8)$$

For instance, using the same $\pi = \text{year}$ and $k = 1$ as above, we can build the filter $f_{\pi, k} = \text{last_year}$, which selects all events that occurred last year.

As the output of a filter applied to a memory M is another memory object $M' \subseteq M$, we can compose filter functions. A sequence of such filters is called a *retrieval path*

$$p = (f_{\pi_1, k_1}, \dots, f_{\pi_n, k_n}) \quad (9)$$

and by definition $p(M) = f_{\pi_n, k_n}(\dots(f_{\pi_1, k_1}(M)))$. In case the result of the operation $p(M)$ contains a single element e , we say that the path p *retrieves* the element e from M , and write $p(M) = e$. In the example shown in Figure 1, the three filters f_1, f_2, f_3 form a retrieval path retrieving the event “last year’s hottest day” from the base memory M_0 .

B. Description complexity of events

As presented in sec. II, we are interested in computing an approximation of the description complexity of an event e . From the above definitions, if there is a path p retrieving e from the base memory M_0 , i.e. $p(M_0) = e$, this path provides a possible unambiguous description for e . We therefore define the description complexity of e as the minimum complexity of a path p retrieving e from the base memory M_0 .

$$K_d(e) = \min_{p \in P_\infty} \{L(p) | p(M_0) = e\} \quad (10)$$

where the bit-length $L(p)$ of a retrieval path is defined as the number of bits of a string encoding the path. If we limit ourselves to prefix-free strings encoding predicates and arguments, the total bit length is given by:

$$L(p) = L((f_{\pi_1, k_1}, \dots, f_{\pi_n, k_n})) \quad (11)$$

$$= L(\pi_1) + L(k_1) + \dots + L(\pi_n) + L(k_n) \quad (12)$$

By considering only a finite number of possible predicates π and arguments k , and a maximum path length, we can

²In the original paper [3], unusually complex events could be handled by considering complexity itself as a way to describe the event: see “the Pisa Tower effect” [8]

construct a finite set P of possible retrieval paths. By limiting the search over this set, we can upper bound the description complexity, and use this upper bound as an approximation:

$$K_d(e) \leq \min_{p \in P \wedge p(M_0)=e} L(p) = \min_{p \in P \wedge p(M_0)=e} \sum_{f_{\pi,k} \in p} L(\pi) + L(k) \quad (13)$$

Algorithm 1: Iterative computation of the approximate complexity

```

1 current_explore ← [(M, 0)] ;
2 future_explore ← [] ;
3 pass ← 0 ;
4 K(e) ← +∞ ;
5 while current_explore ≠ [] and pass < max_pass do
6   for (M_prev, K_prev) ∈ current_explore do
7     for π ∈ P do
8       for k ∈ {0, 1}* do
9         K_current ← l(π) + l(k) + K_prev ;
10        if K_current > max_complex then
11          break ;
12        end
13        M' ← f_{π,k}(M_prev) ;
14        if M' = {e} then
15          K(e) ← min(K(e), K_current) ;
16        else
17          future_explore.append((M', K_current));
18        end
19      end
20    end
21  end
22  current_explore ← future_explore ;
23  future_explore ← [] ;
24  pass ← pass + 1;
25 end

```

The approximation of description complexity from Equation 13 allows for a direct implementation, which is shown in Algorithm 1. This algorithm operates iteratively: starting with the base memory M_0 (line 1), we apply all possible predicate concepts π from a given finite set Π and programs k (lines 6-7), up to a given length `max_len` bits, and apply them: $M' = f_{\pi,k}(M)$ (line 12). We then store the pairs $(M', \text{len}(\pi, k))$ in an array `future_explore`. At the end of the iteration, the results of the filters become the memories which will be explored during the next iteration (lines 21–23). Each pass thus explore retrieval paths of increasing length. When a singleton memory is reached, the complexity of its unique element is upper-bounded with the length of the corresponding retrieval path (line 14).

C. Computing Memorability

As stated in Equation 5, we define memorability $M(e)$ as the absolute difference between the description complexity of

an event and its expected value. As we've just defined $K_d(e)$ and provided an approximation in Equation 13, we now focus on defining the description expected complexity of an event, $K_{exp}(e)$.

As intended in the Equation 5, this term evaluates the complexity the user, or the system, would expect the event e to have without actually describing it, based on their previous knowledge. In our canvas, this prior knowledge consists of the base memory M_0 . The expected complexity of the event e can be computed with a simple first-order approximation, i.e. estimating the average complexity of “similar events” over the base memory M_0 .

Still, difficulty remains in the definition of what can be considered *similar* events. Given that we deal with non comparable events, the only possibility to universally define the notion of similarity is to once again refer to *predicates*. Thus, for a given event e , and a given predicate π_k , we define a π_k neighborhood of e as the set $N_{\pi,k}(e)$ of all other events satisfying π_k .

$$N_{\pi,k}(e) = \{e' \in M_0, \pi_k(e') \wedge e' \neq e\} \quad (14)$$

Now, when considering, for all possible predicates π_k , the corresponding neighborhoods $N_{\pi,k}(e)$, with the convention that $N_{\pi,k}(e) = \emptyset$ if e does not satisfy π_k , we can compute an average expected complexity for e :

$$K_{exp}(e) = \frac{\sum_{\pi,k} \sum_{e' \in N_{\pi,k}(e)} K_d(e')}{\sum_{\pi,k} |N_{\pi,k}(e)|} \quad (15)$$

This definition is consistent with the intuitive ideas that more similar events should weight more in the computation. Indeed, if e' is *very* similar to e , it will appear in many neighborhoods, since it satisfies mostly the same predicates as e . Therefore, it will be present in more terms in Equation 15, and will weight more in the final result.

D. From memorability to abduction

The problem of the abductive inference is different from the basic computation of the memorability score. Here, *knowing* that we want to find a cause c to an observed effect e , we try to find the most remarkable event in past memory in that regard. While our “memorability” score aims to identify the most remarkable events, it does not take into account the added knowledge about the effect e that is available for abductive inference.

This added knowledge can be integrated into the description complexity definition by using conditioned predicates. That is considering as given, and therefore “free” in terms of complexity, information contained in the cause effect e . For instance, when looking for a cause for an anomaly in the living-room, other anomalies occurring in the same living-room will be simpler, since the location “living-room” is already known from knowing the consequence.

Formally, we now consider only predicates where knowledge of the effect event e is appended to all programs k : $\pi_{k::e}(c)$, where $::$ is the append operation. The set of paths

obtained with such predicates is P_e^∞ . This append operation is free in terms of bit-length in the computation of complexity, since the effect event e is an input of the problem. Therefore, we have $L'(\pi_{k::e}) = L(\pi_k) = L(\pi) + L(k)$. Which finally gives the definition for the conditional description complexity:

$$K_d(c|e) = \min_{p \in P_e^\infty} \{L'(p), \quad p(M_0) = c\} \quad (16)$$

$$= \min_{p \in P_e^\infty} \left\{ \sum_{f_{\pi, k::e} \in p} L(\pi) + L(k), \quad p(M_0) = c \right\} \quad (17)$$

This new conditional description complexity translates the additional information provided to the system when answering a user's request. It can then be averaged over similar events to compute the expected conditional description complexity, $K_{exp}(c|e)$. From this, we come to the definition of the conditional memorability, which measures how memorable an event c is, considering the knowledge of another event e :

$$M(c|e) = |K_{exp}(c|e) - K_d(c|e)| \quad (18)$$

Conditional memorability encapsulates the idea presented as the motivation of this paper: when confronted to a surprising situation, and with no other source of information, we can identify events that appear more memorable than others, with regards to the target effect event. As such, our conditional memorability score provides a ranking that can be used for abductive inference, proposing as best candidates the most memorable ones.

IV. IMPLEMENTATION EXAMPLE

A. Setup

To test our approach, we set up an experimental smart home setup. This choice of configuration is motivated by the challenges posed by smart homes for abductive inference: i) as the number of connected devices increases, more events are recorded, making the detection of memorable events more important; ii) smart homes are prone to undergo atypical situations, highly dependent on the context, for which pre-established relations might fail to find good abduction candidates. Furthermore, the choice was also motivated by the ease of application: as previous works exist on smart homes, it is possible to simulate their behavior and quickly generate data to extract events from and test our methods.

To carry out the simulation, we built custom modules into the existing iCasa smart home simulator [9]. This simulation platform allows for simulating autonomic systems, handling internal communications, injection of new components at runtime, deletion or change of existing components. As such, we used a basic scenario consisting of a house of 4 rooms, a single user, and an outdoor zone. All four rooms of the house are equipped with a temperature controller system, monitoring and controlling heaters (fig. 2).

Using this basis, we implemented a scenario spanning over 420 days, and comprising a daily cycle of outdoor weather



Fig. 2. View of the simulator's web interface provided by iCasa. The four rooms are visible, with their equipment and the user.

(temperature and sunlight), as well as user's movements. All these daily changes create non-noticeable events, serving as a background noise for our experiments. To produce outstanding events, we randomly generated around twenty events, spanning over the whole duration of the simulation, of different kinds:

- Unusual weather: the outdoor conditions are set to unusually high or low temperatures.
- Heater failures: heater can fail, making them turning off regardless of the command they receive.
- User's vacation: the user go out of the building for an extended period of time.
- Device removal/addition: a device is removed, or another one is added to the system.




Fig. 3. Time series data from the simulation: outdoor temperature (blue) and controller temperature of a room (orange). On different occasion, the controlled temperature deviates from its setting (21°C). We will evaluate if our system finds these deviations as memorable, and if it can propose causal hypotheses.

The values of all devices and zones variables was regularly monitored throughout the simulation run, and the resulting data, which an excerpt is shown in figure 3 can later be used as a basis for our experiments.

B. Implementing the complexity computation

For the implementation of our method, we first needed to identify and characterize events from the time series data generated by the iCasa simulation. Since this is not the focus point of our present work (see sec. V), we simply apply threshold and pre-computed conditions based detection to create a set of events.

This set of events constitutes the basis of the initial memory M_0 used for computations. We then implemented a small number of predicate concepts:

- $\text{label}(e, k)$: whether the event e has the label ranked k^{th} in the memory.
- $\text{rank}(e, a_1, a_2)$: whether the event e has the rank a_1 along axis a_2 .
- $\text{day}(e, k)$: whether the event e occurred k days ago.
- $\text{month}(e, k)$: whether the event e occurred k months ago.
- $\text{location}(e, k)$: whether the event e occurred in zone k .

With a straightforward implementation of memory, predicates and filters, the implementation of alg. 1 worked, but, as expected, took too long to be usable in realistic scenarios with hundreds or thousands of events to consider. In order to facilitate and speed up computations, we also implemented the following improvements:

- The memory object was augmented with various built-in rankings, allowing for faster operations during future filtering. For instance, since the memory object keeps a mapping from timestamps to events, which allows to quickly filter by date without having to loop over each stored element. This mapping, however, is not directly used to retrieve an event from the outside, as to preserve the theoretical model of memory as an unordered set presented in section III.
- Each of these predicates holds the property that, in addition to True and False, they can return another value, None, which is theoretically treated as False but carries the additional information that this predicate concept will also be false for any other element of the memory for any subsequent program k . This allows to effectively break the innermost loop in alg. 1.
- Some of the filters, for instance date or rank filters, were hard-written to select from the pre-computed mappings of the memory objects rather than testing a predicate over all memory elements.

C. Results

1) *Memorability*: The results of the description complexity evaluation, for the described setup, are shown in fig. 4. The entire computation, over 4 iterations (meaning that retrieval paths contained at most 4 filters), took around 30 seconds on a commercial laptop with an i7-8700u CPU.

A main sequence of “usual” events, which complexity is roughly a logarithm of the elapsed time since their occurrence, is visible. This corresponds to events for which the best retrieval path consists of a time description (e.g. “2 months and 12 days ago”). On the other hand, some events stand

out in terms of complexity: some appear simpler, as they can be distinguished by using their rank alongside an axis (“the hottest day”, “the second longest user’s absence”), or the rare occurrence of their kind (“the only fault on the heater”).




Fig. 4. The computed complexities of events with at most retrieval paths of length at most 4. Events of type “day” (blue), “hot” (green), “cold” (red), “device removal”(cyan) are shown.

Computing the “memorability score”, which is shown in fig. 5, highlights these aforementioned events from the main “usual” sequence. Since the computation of this measure treats unusually complex or simple events the same way (from the absolute value operation in Equation 4), we observe that some events are memorable due to their context only. For instance, a temperature anomaly occurring simultaneously to many other anomalies is notable, as it is costlier than expected to distinguish it from its neighbors.

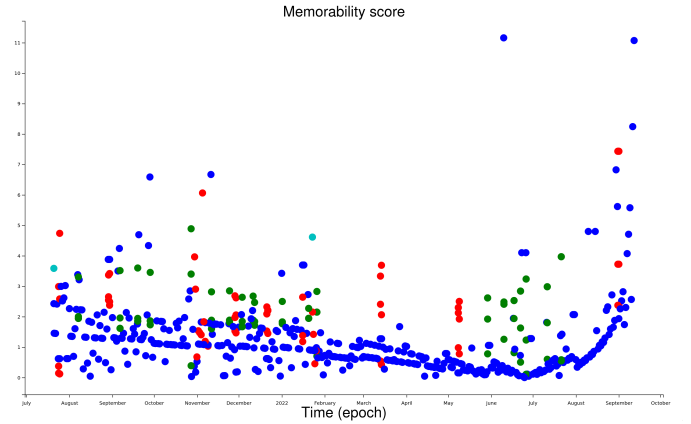


Fig. 5. Memorability score for events in the memory

Given that we generated the data used for this experiment, it is possible to flag all perturbation events from the usual daily events and evaluate how a detection based on “memorability” score would succeed in distinguishing these events. The result is presented as a ROC curve in fig. 6.

2) *Abduction*: As an illustration of the abductive inference possible by using the memorability score, we used as a target




Fig. 6. Experimental ROC curve (True Positive Rate against False Positive Rate) for a classifier based on our memorability score. Measures consider 23 manually flagged events as memorable (events added to the background noise as described in Section IV.)

consequence the temperature drop visible in the raw time series data from our scenario (fig 3).

TODO: the experiment still needs to be done: some additional data is required here (creating the scenario and finding causes.)

3) Discussion:

V. RELATED WORKS

TODO: this section is still WIP: I've added most references and short comments, but I still need to write the story linking them. Our work is intended to be integrated into larger-scale frameworks to monitor and detect events in complex environment such as smart homes. In these works, the approach to smart homes is often regarded as self-organizing systems [10], [11]. As such, they present capacities of adaptation to new goals, new components, new environment. A commonly used approach is the principle of autonomic system, which minimizes user's intervention for management of the system [11], [12].

We want to inscribe our work among other classical concurrent approaches to abduction. In situations where more data is available, we could for instance rely on correlation or causal inference from known relations [13], [14]. Previous relations between inference and complexity have been studied. In fact, the case of inference was one of the motivations for R. Solomonoff to introduce his universal algorithmic probability [15] as a tool to reach an idealized inference machine, creating the notion of complexity concomitantly to Kolmogorov. Subsequently, notions of complexity re-emerged in causal inference: [16] found, when a causal link exists between two random variables, considering the direct joint probability is simpler, in terms of Kolmogorov complexity, than the inverse direction.

More recently, [17] used Minimum Description Length to determine, given a joint probability distribution over (X, Y) ,

whether X causes Y or Y causes X . Their method is based on tree models, and implying that a model respecting the causal relation will be simpler to describe.

[18] PACK algorithm to cluster data using the simplest possible trees: maybe link this to isolation forests?

The topic of finding events from streams of time series data has already been explored in many ways. [19] provides good review of modern approaches and techniques in the field. Some previous work can also be noted for having used AIT techniques to qualify and detect events in time series data. For instance, [20], [21] propose weighted permutation entropy as a proxy for complexity measures in time series data, and use it to find relations between different time series. [22] proposes a MDL approach to find the intrinsic dimensions of time series. All these approaches are interesting, and can be integrated into our canvas as tools to detect events using only complexity. As such, one could achieve a purely complexity-driven process for detecting and qualifying events.

Among various methods, the approach of “Isolation forests” [23], [24] is closely related, by design, to ours. It evaluates the isolation of data points by constructing random binary tree classifiers. On average, outlier points will require less operations to be singled out. Using the average height of leaves in the tree as a metrics, this approach succeeds in identifying outlier points without having to define a “typical” point. This approach can be understood in terms of complexity: each node of binary tree classifier needs a fixed amount of information to be described (which variable and threshold are used). So nodes that place higher need less information to be described. As such, outliers need less information to be singled out. Compared to ours, this method is tailored for data-points living in the same metric space. By using predicates as a proxy for complexity computation, our methods is more general, as it is agnostic regarding the nature of events. However, the introduction of predicates adds a subjectivity in the determination of memorable points, as we will discuss later.

While all these works advocate in favor of a strong link between complexity and the discovery of causes, not other works extends the notion up to the point we propose in this paper, using complexity only as a tool to express the intuitive notion of memorability, and using it for inference.

VI. PERSPECTIVES

The main purpose of the present paper is to show the possible connections between existing definitions of simplicity from cognitive science, Algorithmic Information Theory and a practical use case in cyber-physical systems. As such, many further improvements can be done to pave the way towards a better integration and performance for anomaly detection or abduction.

First, the main limitation of the current approach is the requirement of predefined predicate concepts, from which the different filters are constructed. As an extension, we suggest that in the future, we could explore online generation of such predicate. A possibility would be to analyze discriminating

dimensions of incoming data and create predicate as to name these differences, similar to the contrast operations proposed in [25], [26]. For instance, the predicate concept *hot* can be discovered by discriminating a recent hot day along the temperature axis and naming the difference with the prototypical day.

While the execution time is not part of the theoretical view of complexity, it is of prime importance for practical applications, especially when one considers implementation into real-time systems or embedded devices. While the computation we propose appear to be heavy, and possibly heavier as the number of allowed predicates grows, significant time savings can be achieved by trimming the base memory of past events deemed the most “non memorable”. For instance, one can only retain the 100 most memorable events from the past. The difficulty with this approach is that such operations should be done in a manner to not interfere with the complexity computations for new elements: by forgetting some past events, even uninteresting ones, one should make sure to keep track of what made the interesting ones, interesting. Investigation of how to do so can pave the way towards practical implementations and dynamic selection of interesting events and help reducing the memory and computation cost of data-driven applications.

VII. CONCLUSION

REFERENCES

- [1] L. Magnani, *Abduction, reason and science: Processes of discovery and explanation*. Springer Science & Business Media, 2011.
- [2] M. Li, P. Vitányi, and others, *An introduction to Kolmogorov complexity and its applications*. Springer, 2008, vol. 3.
- [3] J.-L. Dessalles, “Coincidences and the encounter problem: A formal account,” *arXiv preprint arXiv:1106.3932*, 2011.
- [4] J.-P. Delahaye and H. Zenil, “Numerical evaluation of algorithmic complexity for short strings: A glance into the innermost structure of randomness,” *Applied Mathematics and Computation*, vol. 219, no. 1, pp. 63–77, 2012.
- [5] A. N. Kolmogorov, “Three approaches to the definition of the concept “quantity of information,”” *Problemy peredachi informatsii*, vol. 1, no. 1, pp. 3–11, 1965, publisher: Russian Academy of Sciences, Branch of Informatics, Computer Equipment and ...
- [6] P.-A. Murena, M. Al-Ghossein, J.-L. Dessalles, A. Cornuéjols, and others, “Solving Analogies on Words based on Minimal Complexity Transformations,” in *International Joint Conference on Artificial Intelligence*. IJCAI, 2020.
- [7] N. Chater and P. Vitányi, “Simplicity: A unifying principle in cognitive science?” *Trends in cognitive sciences*, vol. 7, no. 1, pp. 19–22, 2003, publisher: Elsevier.
- [8] J.-L. Dessalles, “The Pisa Tower effect.” [Online]. Available: <https://simplicitytheory.telecom-paris.fr/Pisa.html>
- [9] P. Lalanda, E. Gerber-Gaillard, and S. Chollet, “Self-Aware Context in Smart Home Pervasive Platforms,” in *2017 IEEE International Conference on Autonomic Computing (ICAC)*, 2017, pp. 119–124.
- [10] J. Kramer and J. Magee, “A rigorous architectural approach to adaptive software engineering,” *Journal of Computer Science and Technology*, vol. 24, no. 2, pp. 183–188, 2009.
- [11] S. Kounev, P. Lewis, K. L. Bellman, N. Bencomo, J. Camara, A. Diaconescu, L. Esterle, K. Geihs, H. Giese, S. Götz, and others, “The notion of self-aware computing,” in *Self-Aware Computing Systems*. Springer, 2017, pp. 3–16.
- [12] J. O. Kephart and D. M. Chess, “The vision of autonomic computing,” *Computer*, vol. 36, no. 1, pp. 41–50, 2003, publisher: IEEE.
- [13] J. Peters, D. Janzing, and B. Schölkopf, *Elements of causal inference: foundations and learning algorithms*. MIT press, 2017.
- [14] K. Fadiga, E. Houzé, A. Diaconescu, and J.-L. Dessalles, “To do or not to do: finding causal relations in smart homes,” in *202 IEEE International Conference on Autonomic Computing and Self-Organizing Systems (ACSOS)*, 2021, eprint: 2105.10058.
- [15] R. J. Solomonoff, “A formal theory of inductive inference. Part I,” *Information and control*, vol. 7, no. 1, pp. 1–22, 1964, publisher: Elsevier.
- [16] D. Janzing and B. Schölkopf, “Causal inference using the algorithmic Markov condition,” *IEEE Transactions on Information Theory*, vol. 56, no. 10, pp. 5168–5194, 2010, publisher: IEEE.
- [17] A. Marx and J. Vreeken, “Causal inference on multivariate and mixed-type data,” in *Joint European Conference on Machine Learning and Knowledge Discovery in Databases*. Springer, 2018, pp. 655–671.
- [18] N. Tatti and J. Vreeken, “Finding good itemsets by packing data,” in *2008 Eighth IEEE International Conference on Data Mining*. IEEE, 2008, pp. 588–597.
- [19] C. C. Aggarwal, *Outlier Analysis*. Springer International Publishing, 2017.
- [20] G. E. Batista, X. Wang, and E. J. Keogh, “A complexity-invariant distance measure for time series,” in *Proceedings of the 2011 SIAM international conference on data mining*. SIAM, 2011, pp. 699–710.
- [21] B. Fadlallah, B. Chen, A. Keil, and J. Príncipe, “Weighted-permutation entropy: A complexity measure for time series incorporating amplitude information,” *Physical Review E*, vol. 87, no. 2, p. 022911, 2013.
- [22] B. Hu, T. Rakthanmanon, Y. Hao, S. Evans, S. Lonardi, and E. Keogh, “Discovering the intrinsic cardinality and dimensionality of time series using MDL,” in *2011 IEEE 11th International Conference on Data Mining*. IEEE, 2011, pp. 1086–1091.
- [23] F. T. Liu, K. M. Ting, and Z.-H. Zhou, “Isolation forest,” in *2008 Eighth IEEE International Conference on Data Mining*, 2008, pp. 413–422.
- [24] S. Hariri, M. C. Kind, and R. J. Brunner, “Extended Isolation Forest,” *IEEE Transactions on Knowledge and Data Engineering*, vol. 33, no. 4, pp. 1479–1489, 2021.
- [25] J.-L. Dessalles, “From conceptual spaces to predicates,” in *Applications of conceptual spaces: The case for geometric knowledge representation*, F. Zenker and P. Gärdenfors, Eds. Dordrecht: Springer, 2015, pp. 17–31. [Online]. Available: http://www.dessalles.fr/papers/Dessalles_14050201.pdf
- [26] P. Gärdenfors, *Conceptual spaces: The geometry of thought*. MIT press, 2004.